

Student Name: _____

Student Identification Number: _____

Course Abbreviation and Number: ECE459
 Course Title: Programming for Performance
 Section(s): 001
 Sections Combined Course(s):
 Section Numbers of Combined Course(s):
 Instructor(s): Patrick Lam
 Date of Exam: April 18, 2011
 Exam Period Start time: 9:00 End time: 11:30
 Duration of Exam: 2.5 hours
 Number of Exam Pages: (includes cover page) 5
 Exam Type: (select one) ☐ Closed Book ☐ Special Materials ☒ Open Book

Materials Allowed: (select one)

☐ No additional materials are allowed.

☒ Materials allowed are listed below:

Any printed materials. No electronic devices.

Exams are printed double sided on white paper.

☐ Select this box if second side of paper is to be used for rough work calculations.

Marking Scheme:

Question	Score	Question	Score
1		6	
2		7	
3		8	
4		9	
5		10	

Programming for Performance (ECE459): Final

April 18, 2010

This open-book final has 6 questions, worth 20 points each. Answer all questions. Write the answers in your answer book. You may consult any printed material (books, notes, etc).

Question 1: Short Answer

One point each. **Answer these questions in your exam booklet.**

- (a) Two library functions for doing communication in MPI are _____ and _____.
- (b) Doing more work in parallel increases a system's _____.
- (c) Modern CPUs spend most of their time waiting due to _____.
- (d) When using OpenCL, you may choose to divide the computation space into _____.
- (e) Gustafson's Law says that parallelization isn't hopeless when you can increase the _____.
- (f) One condition that would impede automatic parallelization is _____.
- (g) One primary design goal of DTrace was _____.
- (h) A different primary design goal of DTrace was _____.
- (i) An **sfence** instruction prevents reordering of _____.
- (j) To effectively parallelize an OpenMP loop where different iterations run for different amounts of time, you want to use _____.
- (k) **oprofile** is an example of a _____ profiler.
- (l) The three steps in using profile-guided optimization are: _____.
- (m) The main difference between combine and reduce in MapReduce is that _____.
- (n) The obvious way to get rid of a race condition is by using _____.
- (o) A technique to get rid of a WAR dependency is _____.
- (p) The first thing you should do before trying to improve performance of some code is to _____ it.

- (q) The term for speeding things up by doing many things at once is _____.
- (r) Re-entrancy for a lock means that you can _____.
- (s) Botnets are a good example of this parallelization pattern: _____.
- (t) The bit of code that runs massively parallel in GPU programming is a _____.

Question 2: OpenMP

Consider the following code¹.

```

1 # define NV 4
2
3 /* don't worry about mind, connected */
4 # pragma omp parallel /* private, shared etc */
5 {
6     my_id = omp_get_thread_num ( );
7     nth = omp_get_num_threads ( );
8     my_first = ( my_id * NV ) / nth;
9     my_last = ( ( my_id + 1 ) * NV ) / nth - 1;
10    # pragma omp single
11    {
12        printf ( "P%d: Parallel region begins with %d threads\n",
13                my_id, nth);
14        printf ( "\n");
15    }
16    fprintf ( stdout, "P%d: First=%d Last=%d\n",
17             my_id, my_first, my_last);
18
19    for ( my_step = 1; my_step < NV; my_step++ )
20    {
21        # pragma omp single
22        {
23            md = i4_huge;
24            mv = -1;
25        }
26        find_nearest ( my_first, my_last, mind, connected, &my_md, &my_mv);
27        # pragma omp barrier
28    }
29 }
```

- (a) Explain what each of the OpenMP pragmas does.
- (b) Assume that OMP_NUM_THREADS is 4 and draw a diagram explaining what the threads do.

¹http://people.sc.fsu.edu/~jburkardt/c_src/dijkstra_open_mp/dijkstra_open_mp.html.

Question 3: Reductions

We saw this example of a reduction in the notes.

```
1 double sum (double *array, int length)
2 {
3     double total = 0;
4
5     for (int i = 0; i < length; i++)
6         total += array[i];
7     return total;
8 }
```

I mentioned that the Solaris compiler could detect that it was a reduction and parallelize it.

- (a) Write down, in reasonably-detailed pseudocode, the corresponding parallelized code.
- (b) Write down the assumptions that you're making about the + operator in your parallelization.

Question 4: Memory Barriers and Consistency Models

Consider the following code; all variables are initially 0.

```
T1: x = 1; r1 = y;
T2: y = x; r2 = x;
```

Assume the architecture is not sequentially consistent.

- Show me all possible (intermediate and final) memory values and how they arise.

Question 5:

Here is some C code from meschach.

```
1 double zm_norm1(ZMAT *A)
2 {
3     int i, j, m, n;
4     Real maxval, sum;
5
6     if (A == ZMNULL)
7         error(E_NULL, "zm_norm1");
8
9     m = A->m; n = A->n;
10    maxval = 0.0;
```

```

11
12     for ( j = 0; j < n; j++)
13     {
14         sum = 0.0;
15         for ( i = 0; i < m; i ++ )
16             sum += zabs(A->me[ i ][ j ] );
17         maxval = max(maxval, sum);
18     }
19
20     return maxval;
21 }

```

- Describe 2 compiler optimizations that could apply to the code and the resulting code after optimization. Briefly summarize the conditions that need to hold for the optimizations to be safe; you may add qualifiers to the code if you want.

Question 6: GPU Programming

Here is some C code².

```

# define M 500
# define N 500

int i, j;
double diff;
double u[M][N], w[M][N];

diff = 0.0;
for ( i = 1; i < M - 1; i++ ) {
    for ( j = 1; j < N - 1; j++ ) {
        w[i][j] = ( u[i-1][j] + u[i+1][j] + u[i][j-1] + u[i][j+1] ) / 4.0;

        if ( diff < fabs ( w[i][j] - u[i][j] ) )
            diff = fabs ( w[i][j] - u[i][j] );
    }
}

```

- Express the code as one or many OpenCL kernels. (I don't care about host code, just the kernel itself.)
- Indicate the floating-point operations in the kernel.
- Do you need to worry about synchronization?

²http://people.sc.fsu.edu/~jburkardt/c_src/heated_plate/heated_plate.c