

# Programming for Performance (ECE459): Midterm

## February 26, 2014

This open-book midterm has 3 pages and 4 questions, worth 25 points each. Answer the questions in your answer book. You may consult any printed material (books, notes, etc).

### 1 Short Answer

Answer these questions using at most three sentences. Each question is worth 2.5 points.

- (a) What can the compiler assume about a **restrict**-qualified pointer?
- (b) Write down a Write-after-Write dependency. Rewrite your code, eliminating the write-after-write dependency. (You may, of course, introduce a different dependency).
- (c) You are running a simple web server on an otherwise-unloaded 8-core machine. The web server works as follows: when a main thread accepts a connection, it dispatches a thread from a thread pool to respond to the request. Do you expect better throughput from a pool with 8 or 9 threads? Why? (Be explicit with your assumptions.)
- (d) OpenMP will not parallelize this loop properly. Propose an equivalent for loop which will parallelize.

```
1  double * array = malloc(sizeof(double) * 20);
2  for (double d = 0.0; d < 10.0; d += 0.5) {
3      array[(int)(d*2)] = sin(d);
4  }
```

- (e) Will you ever get a race condition from converting an OpenMP shared variable into a private variable? Why or why not?
- (f) Gustafson's Law differs from Amdahl's Law because it allows what to vary?
- (g) What is one problem with keeping a bunch of joinable threads around indefinitely?
- (h) Say you have 300,000 potentially-active incoming connections open, but only 5 of them are ever active at once. Would threads or nonblocking I/O be better? Why?
- (i) Which parallelization pattern most closely corresponds to a bank of subway turnstiles all controlling access to the subway in parallel?
- (j) Give an example where you would use OpenMP tasks rather than sections. Explain why sections don't work in that case.

## 2 Locking

Louis Reasoner is working on the following tree implementation.

```
1 struct node
2 {
3     struct node * left , * right;
4     int key;
5     int * data;
6 };
7
8 struct node * root;
9
10 int find_and_increment(int key)
11 {
12     pthread_mutex_t global_lock = PTHREAD_MUTEX_INITIALIZER;
13
14     struct node * n = root;
15     pthread_mutex_lock(&global_lock);
16     while (n != NULL) {
17         if (key == n->key) {
18             *n->data++;
19             pthread_mutex_unlock(&global_lock);
20             return *n->data;
21         }
22         if (key < n->key)
23             n = n->left;
24         else if (key > n->key)
25             n = n->right;
26     }
27     pthread_mutex_unlock(&global_lock);
28     return NULL;
29 }
```

(a, 5 points) This code does not actually lock accesses to the tree. Why not? Propose a fix which properly locks accesses to the tree.

(b, 20 points) Make the following assumptions: (i) the **data** pointers may be shared among nodes and may be changed (as we see in the example); (ii) the structure of the tree (**key**, **left** and **right** fields) never changes after the tree is initialized. Now, propose changes to **struct node** and **find\_and\_increment** which permit two threads to concurrently call the function, while avoiding races on the **data** fields. Explain why your changes are correct.

### 3 Reductions

If you ask a compiler to parallelize the following loop, it will tell you that it found a reduction.

```
1  double sum(double[] array, int N) {
2      double accum = 0.0;
3      for (int i = 0; i < N; i++)
4          accum += array[i];
5      return accum;
6  }
```

Assume that there is a `NUM_THREADS` constant. You have to use that number of threads (part a) or tasks (part b). For simplicity, assume that `N % NUM_THREADS == 0`.

(a, 10 points) Rewrite this loop using `pthread` primitives to implement the reduction. (Casting `int` to and from `void *` is OK here.)

(b, 10 points) Rewrite this loop using OpenMP directives (no reduction).

(c, 5 points) Describe the source of overheads for autoparallelized reduction and one of pthreads/OpenMP. As the array gets larger, which implementation is fastest?

### 4 Dependences and parallelization

Consider these two code fragments.

```
1  void methodA(double * x, double * y) {
2      for (int i = 0; i < 101; ++i)
3          for (int j = 0; j < 100; ++j)
4              x[j] = y[i+1];
5  }

1  void methodB(double * x) {
2      for (int i = 0; i < 1000; ++i)
3          for (int j = 0; j < 100; ++j)
4              x[j] += j;
5  }
```

For each of these code fragments:

(a, 2.5 points) Describe the dependency in the code fragment. Be specific.

(b, 10 points) Will a compiler auto-parallelize this loop fragment? Why or why not? If the answer is “no”, propose transformations which will allow the loop fragment to be auto-parallelized.