

# **Lecture 02—Amdahl's Law, Modern Hardware**

## **ECE 459: Programming for Performance**

Patrick Lam

University of Waterloo

January 7, 2015

# About Prediction and Speedups

Cliff Click said: “5% miss rates dominate performance.”

Why is that?

# About Prediction and Speedups

Cliff Click said: “5% miss rates dominate performance.”

Why is that?

Recall: 100-1000 slot penalty for a miss.

See `L02.pdf` for a calculation.

# Forcing Branch Mispredicts

blog.man7.org/2012/10/  
how-much-do-builtinexpect-likely-and.html

```
#include <stdlib.h>
#include <stdio.h>

static __attribute__((noinline)) int f(int a) { return a; }

#define BSIZE 1000000
int main(int argc, char* argv[])
{
    int *p = calloc(BSIZE, sizeof(int));
    int j, k, m1 = 0, m2 = 0;
    for (j = 0; j < 1000; j++) {
        for (k = 0; k < BSIZE; k++) {
            if (__builtin_expect(p[k], EXPECT_RESULT)) {
                m1 = f(++m1);
            } else {
                m2 = f(++m2);
            }
        }
    }

    printf("%d, %d\n", m1, m2);
}
```

Running times: 3.1s with good (or no) hint, 4.9s with bogus hint.

# Limitations of Speedups

Our main focus is parallelization.

- Most programs have a sequential part and a parallel part; and,
- Amdahl's Law answers, “what are the limits to parallelization?”

## Formulation (1)

$S$ : fraction of serial runtime in a serial execution.

$P$ : fraction of parallel runtime in a serial execution.

Therefore,  $S + P = 1$ .

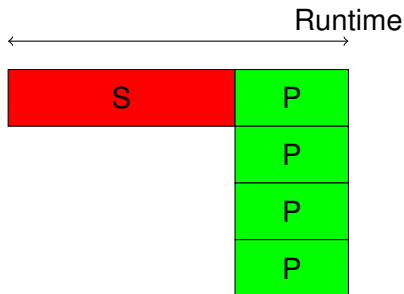
With 4 processors, best case, what can happen to the following runtime?



# Formulation (1)



We want to split up the parallel part over 4 processors



## Formulation (2)

$T_s$ : time for the program to run in serial

$N$ : number of processors/parallel executions

$T_p$ : time for the program to run in parallel

- Under perfect conditions, get  $N$  speedup for  $P$

$$T_p = T_s \cdot \left( S + \frac{P}{N} \right)$$




## Formulation (3)

How much faster can we make the program?

$$\begin{aligned} speedup &= \frac{T_s}{T_p} \\ &= \frac{T_s}{T_s \cdot (S + \frac{P}{N})} \\ &= \frac{1}{S + \frac{P}{N}} \end{aligned}$$

(assuming no overhead for parallelizing; or costs near zero)

# Fixed-Size Problem Scaling, Varying Fraction of Parallel Code



# Amdahl's Law

Replace  $S$  with  $(1 - P)$ :

$$\textit{speedup} = \frac{1}{(1-P) + \frac{P}{N}}$$

$$\textit{maximum speedup} = \frac{1}{(1-P)}, \text{ since } \frac{P}{N} \rightarrow 0$$

As you might imagine, the asymptotes in the previous graph are bounded by the maximum speedup.

# Assumptions behind Amdahl's Law

How can we invalidate Amdahl's Law?

# Assumptions behind Amdahl's Law

We assume:

- problem size is fixed (we'll see this soon);
- program/algorithm behaves the same on 1 processor and on  $N$  processors; and
- that we can accurately measure runtimes—  
i.e. that overheads don't matter.