

Lecture 05—Processes vs Threads II; Thread Usage

ECE 459: Programming for Performance

Patrick Lam

University of Waterloo

January 12, 2015

Processes: fork Usage Example (OS refresher)

```
pid = fork();  
if (pid < 0) {  
    fork_error_function();  
} else if (pid == 0) {  
    child_function();  
} else {  
    parent_function();  
}
```

fork produces a second copy of the calling process, which starts execution after the call.

The only difference between the copies is the return value: the parent gets the pid of the child, while the child gets 0.

Live coding: overheads of threads vs fork

Results: Threads Offer a Speedup of 6.5 over fork

Here's a benchmark between fork and Pthreads on a laptop, creating and destroying 50,000 threads:

```
jon@riker examples master % time ./create_fork
0.18s user 4.14s system 34% cpu 12.484 total
jon@riker examples master % time ./create_pthread
0.73s user 1.29s system 107% cpu 1.887 total
```

Clearly Pthreads incur much lower overhead than fork.

Assumptions

First, we'll see how to use threads on “embarrassingly parallel problems”.

- mostly-independent sub-problems (little synchronization); and
- strong locality (little communication).

Later, we'll see:

- which problems can be parallelized ([dependencies](#))
- alternative parallelization patterns
(right now, just use one thread per sub-problem)

POSIX Threads

- Available on most systems
- Windows has Pthreads Win32, but I wouldn't use it; use Linux for this course
- API available by `#include <pthread.h>`
- Compile with pthread flag
(`gcc -pthread prog.c -o prog`)

C++ 11 Threads

- Now part of the C++ standard (library)
- API available with `#include <thread>`
- Compile with flags:
(`g++ -std=c++11 -pthread prog.c -o prog`)

Pthreads: Creating Threads

```
int pthread_create(pthread_t* thread ,
                  const pthread_attr_t* attr ,
                  void* (*start_routine)(void*) ,
                  void* arg );
```

thread: creates a handle to a thread at pointer location

attr: thread attributes (NULL for defaults, more details later)

start_routine: function to start execution

arg: value to pass to start_routine

returns 0 on success, error number otherwise
(contents of *thread are undefined)

Creating Threads—Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    printf("In main\n");
}
```

Simply creates a thread and terminates
(usage isn't really right, as we'll see.)

Creating Threads—C++11 Example

```
#include <thread>
#include <iostream>

void run() {
    std::cout << "In run\n";
}

int main() {
    std::thread t1(run);
    std::cout << "In main\n";
    t1.join(); // hang in there...
}
```

Waiting for Threads

```
int pthread_join(pthread_t thread,  
                 void** retval)
```

thread: wait for this thread to terminate (thread must be joinable).

retval: stores exit status of thread (set by `pthread_exit`) to the location pointed by `*retval`. If cancelled, returns `PTHREAD_CANCELED`. `NULL` is ignored.

returns 0 on success, error number otherwise.

Only call this one time per thread! Multiple calls on the same thread leads to undefined behaviour.

Waiting for Threads—Pthreads example

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    printf("In main\n");
    pthread_join(thread, NULL);
}
```

This now waits for the newly created thread to terminate.

Creating Threads—C++11 Example

```
#include <thread>
#include <iostream>

void run() {
    std::cout << "In run\n";
}

int main() {
    std::thread t1(run);
    std::cout << "In main\n";
    t1.join(); // aha!
}
```

Passing Data to Pthreads threads... Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)&i);
```

This is a **terrible** idea. Why?

Passing Data to Pthreads threads... Wrongly

Consider this snippet:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)&i);
```

This is a **terrible** idea. Why?

- ❶ The value of `i` will probably change before the thread executes
- ❷ The memory for `i` may be out of scope, and therefore invalid by the time the thread executes

Passing Data to Pthreads threads

What about:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)i);  
  
...  
  
void* run(void* arg) {  
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

Passing Data to Pthreads threads

What about:

```
int i;  
for (i = 0; i < 10; ++i)  
    pthread_create(&thread[i], NULL, &run, (void*)i);  
  
...  
  
void* run(void* arg) {  
    int id = (int)arg;
```

This is suggested in the book, but should carry a warning:

- Beware size mismatches between arguments: no guarantee that a pointer is the same size as an int, so your data may overflow.
- Sizes of data types change between systems. For maximum portability, just use pointers you got from malloc.

Passing Data to C++11 threads

It's easier to get data to threads in C++11:

```
#include <thread>
#include <iostream>

void run(int i) {
    std::cout << "In run " << i << "\n";
}

int main() {
    for (int i = 0; i < 10; ++i) {
        std::thread t1(run, i);
        t1.detach(); // see the next slide ...
    }
}
```

Getting Data from C++11 threads

...but it's harder to get data back.

Use async and future abstractions:

```
#include <thread>
#include <iostream>
#include <future>

int run() {
    return 42;
}

int main() {
    std::future<int> t1_retval =
        std::async(std::launch::async, run);
    std::cout << t1_retval.get();
}
```

Detached Threads

Joinable threads (the default) wait for someone to call `pthread_join` before they release their resources.

Detached threads release their resources when they terminate, without being joined.

```
int pthread_detach(pthread_t thread);
```

thread: marks the thread as detached

returns 0 on success, error number otherwise.

Calling `pthread_detach` on an already detached thread results in undefined behaviour.

Thread Termination

```
void pthread_exit(void *retval);
```

retval: return value passed to function that calls `pthread_join`

`start_routine` returning is equivalent to calling `pthread_exit` with that return value;

`pthread_exit` is called implicitly when the `start_routine` of a thread returns.

There is no C++11 equivalent.

Attributes

By default, threads are *joinable* on Linux, but a more portable way to know what you're getting is to set thread attributes.

You can change:

- Detached or joinable state
- Scheduling inheritance
- Scheduling policy
- Scheduling parameters
- Scheduling contention scope
- Stack size
- Stack address
- Stack guard (overflow) size

Attributes—Example

```
size_t stacksize;  
pthread_attr_t attributes;  
pthread_attr_init(&attributes);  
pthread_attr_getstacksize(&attributes, &stacksize);  
printf("Stack size = %i\n", stacksize);  
pthread_attr_destroy(&attributes);
```

Running this on a laptop produces:

```
jon@riker examples master % ./stack_size  
Stack size = 8388608
```

Setting a thread state to joinable:

```
pthread_attr_setdetachstate(&attributes,  
                             PTHREAD_CREATE_JOINABLE);
```

Detached Threads: Warning!

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    pthread_detach(thread);
    printf("In main\n");
}
```

When I run it, it just prints "In main", why?

Detached Threads: Solution to Problem

```
#include <pthread.h>
#include <stdio.h>

void* run(void*) {
    printf("In run\n");
}

int main() {
    pthread_t thread;
    pthread_create(&thread, NULL, &run, NULL);
    pthread_detach(thread);
    printf("In main\n");
    pthread_exit(NULL); // This waits for all detached
                        // threads to terminate
}
```

Make the final call `pthread_exit` if you have any detached threads. (There is no C++11 equivalent.)

Threading Challenges

- Be aware of scheduling (you can also set affinity with pthreads on Linux).
- Make sure the libraries you use are **thread-safe**:
 - ▶ Means that the library protects its shared data.
- glibc reentrant functions are also safe: a program can have more than one thread calling these functions concurrently.
- **Example:** `rand_r` versus `rand`.