

# Lecture 06—A1; Race Conditions; More Synchronization; Async I/O

January 21, 2014

# Roadmap

Past: Modern Hardware, Threads

Now: A1 discussion, non-blocking I/O;

Next: Race Conditions, Locking

## Last Time

- Processes vs threads.
- Creating, joining and exiting POSIX threads.  
Remember, they are 1:1 with kernel threads and can run in parallel on multiple CPUs.
- Difference between `joinable` and `detached` threads.

# Part I

## Assignment 1

# Your Task

Re-assemble the picture:



# What I provide

Serial C code:

- uses curl to fetch the image over the network;
- uses libpng to stitch together the image.

Plus, a web API to provide images to you.

# What you hand in

Part 1: pthreads parallelized implementation.

- really easy!
- (also, analyze your speedups in the report.)

Part 2: nonblocking I/O implementation.

- more challenging;
- lecture today will help.

Also Parts 0, 3: analysis & discussion.

# Tour of the code

main loop: while still missing some fragments,

- retrieve a fragment over the network;
- copy bits into our array;

Then, write all the bits in one PNG file.



## Notable bits I: retrieving the file

Here's how I retrieve the file:

```
curl_easy_setopt(curl, CURLOPT_URL, url);  
  
// do curl request; check for errors  
res = curl_easy_perform(curl);
```

But wait! I had to tell curl where to put the file:

```
struct bufdata bd;  
bd.buf = input_buffer;  
curl_easy_setopt(curl, CURLOPT_WRITEFUNCTION, write_cb);  
curl_easy_setopt(curl, CURLOPT_WRITEDATA, &bd);
```

My `write_cb` callback function puts data in `input_buffer` (straightforward `memcpy`-based implementation).

## Notable bits II: parsing the fragments

Bunch of `libpng` magic:

`libpng` wants to put the image data in a `png_bytep * array`, where each element points to a row of pixels.

My `read_png_file` function allocates the data; caller must free.

Then, `paint_destination` fills in the output array, pasting together the fragments.

## Notable bits III: writing the output

Well, not that notable. Symmetric to read.

Note: be sure to free everything! (We'll check.)

## Part (a): using pthreads

You might need to refactor the code to parallelize it well.

Start some threads.

Justify why the threads are not interfering. Time the result.

## Part (b): nonblocking I/O

Main subject of this lecture.

Will be more complicated than using threads!

## Part (b)': JavaScript

As an alternate option, you may use either `node.js` or client-side JavaScript to do the nonblocking I/O.

Let me know if you want to do this. You are on your own, though.

## Part II

### Asynchronous/non-blocking I/O

## Juicy Quotes

**Asynchronous I/O on linux**  
**or: Welcome to hell.**

(mirrored at [compgeom.com/~piyush/teach/4531\\_06/project/hell.html](http://compgeom.com/~piyush/teach/4531_06/project/hell.html))

“Asynchronous I/O, for example, is often infuriating.”  
— Robert Love. *Linux System Programming, 2nd ed*, page 215.



# Why non-blocking I/O?

Consider some I/O:

```
fd = open ( ... );  
read ( ... );  
close ( fd );
```

Not very performant—under what conditions do we lose out?

# Mitigating I/O impact

So far: can use threads to mitigate latency.  
What are the disadvantages?

# Mitigating I/O impact

So far: can use threads to mitigate latency.  
What are the disadvantages?

- race conditions
- overhead/max # of thread limitations

Live coding: forkbomb Patrick's laptop!

(well, threadbomb anyway)

# An Alternative to Threads

Asynchronous/nonblocking I/O.

```
fd = open(..., O_NONBLOCK);  
read(...); // returns instantly!  
close(fd);
```

...



(credit: Yskyflyer, Wikimedia Commons)

## Not Quite So Easy: Live Demo

Doesn't work on files—they're always ready. Only e.g. sockets.

## Other Outstanding Problem with Nonblocking I/O

How do you know when I/O is ready to be queried?

# Other Outstanding Problem with Nonblocking I/O

How do you know when I/O is ready to be queried?

- polling (`select`, `poll`, `epoll`)
- interrupts (signals)