

Lecture 10—C++ atomics; Compilers and You

January 26, 2015

Roadmap

Last Time: Synchronization Mechanisms
This Time: C++ atomics; C Compilers;
Dependencies

Part I

C++ atomics

About C++ atomics

You can use the default `std::memory_order`.
(= sequential consistency)

Don't use relaxed atomics unless you're an expert!

[http://stackoverflow.com/questions/9553591/
c-stdatomic-what-is-stdmemory-order-and-how-to-use-them](http://stackoverflow.com/questions/9553591/c-stdatomic-what-is-stdmemory-order-and-how-to-use-them)

Really, don't use C++ relaxed atomics!



C++ atomics: Key Idea

An *atomic operation* is indivisible.

Other threads see state before or after the operation, nothing in between.

Simplest: `atomic_flag`

Represents a boolean flag.

```
#include <atomic>

atomic_flag f = ATOMIC_FLAG_INIT;
```

Operations on `atomic_flag`

Can clear, and can test-and-set:

```
#include <atomic>

atomic_flag f = ATOMIC_FLAG_INIT;
int foo() {
    f.clear();
    if (f.test_and_set()) {
        // was true
    }
}
```

`test_and_set`: atomically sets to true,
returns previous value.

No assignment (=) operator.

Using more general C++ atomics

Declaring them:

```
#include <atomic>

atomic<int> x;
```

Library's implementation:

- on small types, lock-free operations;
- on large types, mutexes.

What to do with Atomics

Kinds of operations:

- reads
- writes
- read-modify-write (RMW)

Reads and Writes

C++ has syntax to make these all transparent:

```
#include <atomic>
#include <iostream>

std::atomic<int> ai;
int i;

int main() {
    ai = 4;
    i = ai;
    ai = i;
    std::cout << i;
}
```

Can also use `i = ai.load()` and `ai.store(i)`.

Read-Modify-Write (RMW)

Consider `ai++`.

This is really

```
tmp = ai.read(); tmp++; ai.write(tmp);
```

Hardware can do that atomically.

Other RMWs: `+-`, `&=`, etc, compare-and-swap

more info:

[http://preshing.com/20130618/
atomic-vs-non-atomic-operations/](http://preshing.com/20130618/atomic-vs-non-atomic-operations/)

Part II

Making C Compilers Work For You

Three Address Code

- An intermediate code used by compilers for analysis and optimization.
- Statements represent one fundamental operation—we can consider each operation **atomic**.
- Statements have the form:
$$result := operand_1 \operatorname{operator} operand_2$$
- Useful for reasoning about data races, and easier to read than assembly.
(separates out memory reads/writes).

GIMPLE

- GIMPLE is the three address code used by gcc.
- To see the GIMPLE representation of your code use the `-fdump-tree-gimple` flag.
- To see all of the three address code generated by the compiler use `-fdump-tree-all`. You'll probably just be interested in the optimized version.
- Use GIMPLE to reason about your code at a low level without having to read assembly.

Live Coding Demo: GIMPLE

Branch Prediction Hints

As seen earlier in class, gcc allows you to give branch prediction hints by calling this builtin function:

```
long __builtin_expect (long exp, long c)
```

The expected result is that `exp` equals `c`.

Compiler reorders code & tells CPU the prediction.

The restrict Keyword

A new feature of C99: “The restrict type qualifier allows programs to be written so that translators can produce significantly faster executables.”

- To request C99 in gcc, use the `-std=c99` flag.

`restrict` means: you are promising the compiler that the pointer will never **alias** (another pointer will not point to the same data) for the lifetime of the pointer.

Example of restrict (1)

Pointers declared with `restrict` must never point to the same data.

From Wikipedia:

```
void updatePtrs(int* ptrA, int* ptrB, int* val) {  
    *ptrA += *val;  
    *ptrB += *val;  
}
```

Would declaring all these pointers as `restrict` generate better code?

Example of restrict (2)

Let's look at the GIMPLE:

```
1 void updatePtrs(int* ptrA, int* ptrB, int* val) {  
2   D.1609 = *ptrA;  
3   D.1610 = *val;  
4   D.1611 = D.1609 + D.1610;  
5   *ptrA = D.1611;  
6   D.1612 = *ptrB;  
7   D.1610 = *val;  
8   D.1613 = D.1612 + D.1610;  
9   *ptrB = D.1613;  
10 }
```

- Could any operation be left out if all the pointers didn't overlap?

Example of restrict (3)

```
1 void updatePtrs(int* ptrA, int* ptrB, int* val) {  
2   D.1609 = *ptrA;  
3   D.1610 = *val;  
4   D.1611 = D.1609 + D.1610;  
5   *ptrA = D.1611;  
6   D.1612 = *ptrB;  
7   D.1610 = *val;  
8   D.1613 = D.1612 + D.1610;  
9   *ptrB = D.1613;  
10 }
```

- If `ptrA` and `val` are not equal, you don't have to reload the data on **line 7**.
- Otherwise, you would: there might be a call
`updatePtrs(&x, &y, &x);`

Example of restrict (4)

Hence, this markup allows optimization:

```
void updatePtrs(int* restrict ptrA ,  
                int* restrict ptrB ,  
                int* restrict val)
```

Note: you can get the optimization by just declaring ptrA and val as restrict; ptrB isn't needed for this optimization

Summary of restrict

- Use `restrict` whenever you know the pointer will not alias another pointer (also declared `restrict`)

It's hard for the compiler to infer pointer aliasing information; it's easier for you to specify it.

⇒ compiler can better optimize your code (more perf!)

Caveat: don't lie to the compiler, or you will get **undefined behaviour**.

Aside: `restrict` is not the same as `const`. `const` data can still be changed through an alias.