

Lecture 27—Profiler Lies

ECE 459: Programming for Performance

March 13, 2015

Part I

Lies about Calling Context

gprof and KCacheGrind

Who can we trust?

Some profiler results are real.

Other results are interpolated, and perhaps wrong.

Reference: Yossi Kreinin,

<http://www.yosefk.com/blog/how-profilers-lie-the-cases-of-gprof-and-kcachegrind.html>

Running Example

```
void work(int n) {  
    volatile int i=0; //don't optimize away  
    while(i++ < n);  
}  
void easy() { work(1000); }  
void hard() { work(1000*1000*1000); }  
int main() { easy(); hard(); }
```

Running the Running Example

```
[plam@lynch L27]$ gprof ./try gmon.out
```

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
101.30	1.68	1.68	2	840.78	840.78	work
0.00	1.68	0.00	1	0.00	840.78	easy
0.00	1.68	0.00	1	0.00	840.78	hard

That's not right!

easy takes ≈ 0 s, hard takes 1.68s.

What's Wrong?

Need to understand how gprof works.

- **profil()**: asks glibc to record which instruction is currently executing (100×/second).
- **mcount()**: records call graph edges; called by -pg instrumentation.

profil information is statistical;
mcount information is exact.

Those Numbers Again

```
[plam@lynch L27]$ gprof ./try gmon.out
```

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
101.30	1.68	1.68	2	840.78	840.78	work
0.00	1.68	0.00	1	0.00	840.78	easy
0.00	1.68	0.00	1	0.00	840.78	hard

- calls: reliable;
- self seconds: sampled, but OK here;
- total ms/call: interpolated!

total ms/call

gprof sees:

- total of 1.68s in work,
- 1 call to work from easy;
- 1 call to work from hard.

All of these numbers are reliable.

gprof's unreliable conclusion:

easy, hard both cause 840ms of work time.

Wrong: work takes $1000\times$ longer when called from hard!

Where gprof guesses: Call graph edges

- contribution of children to parents;
- total runtime spent in self+children;
- etc.

When are call graph edges right?

Two cases:

- functions with only one caller
(e.g. `f()` only called by `g()`); or,
- functions which always take the same time to complete
(e.g. `rand()`).

What's sketchy:

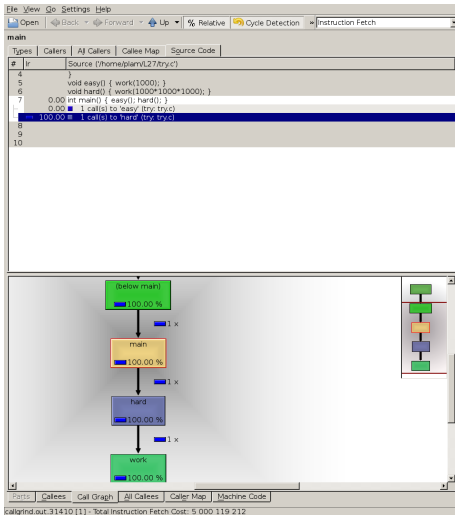
Any function whose running time depends on its inputs,
and which is called from multiple contexts.

KCacheGrind

KCacheGrind is a frontend to callgrind.

callgrind is part of valgrind,
and runs the program under an x86 JIT.

KCacheGrind example



Yes, hard takes all the time.

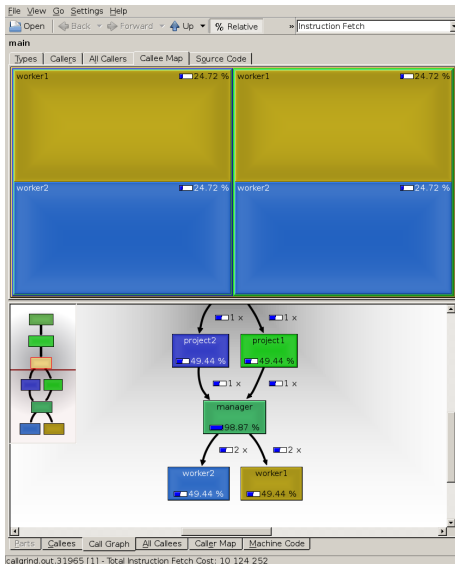
More Complex Example

```
void worker1(int n) {
    volatile int i=0;
    while(i++<n);
}
void worker2(int n) {
    volatile int i=0;
    while(i++<n);
}
void manager(int n1, int n2) {
    worker1(n1);
    worker2(n2);
}
void project1() {
    manager(1000, 1000000);
}
void project2() {
    manager(1000000, 1000);
}
int main() {
    project1();
    project2();
}
```

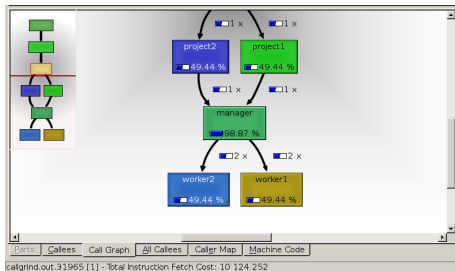
Example explained in 2 lines

Now `worker2` takes all the time in `project1`,
and `worker1` takes all the time in `project2`.

What about KCacheGrind now?



KCacheGrind Truths




This is the call graph.

worker1 and worker2 do each take about 50% of time.
So do project2 and project1.

(gprof would interpolate that too.)

KCacheGrind Lies



KCacheGrind is reporting:

- worker1 and worker2 doing half the work in each project.

That's not what the code says.

Why KCacheGrind Lies

gprof reports time spent in `f()` and `g()`,
and how many times `f()` calls `g()`.

callgrind also reports
time spent in `g()` when called from `f()`,
i.e. some calling-context information.

callgrind does not report time spent in `g()`
when called from `f()` when called from `h()`.

We don't get the `project1` to `manager` to `worker1` link.

(We have Edges but need Edge-Pairs).

gprof/KCacheGrind summary

Some results are exact;
some results are sampled;
some results are interpolated.

If you understand the tool,
you understand where it can go wrong.

Understand your tools!

Part II

Lies from Metrics

Lying perf counters

While app-specific metrics can lie too,
mostly we'll talk about CPU perf counters.

Reference: Paul Khuong,

<http://www.pvk.ca/Blog/2014/10/19/performance-optimisation---writing-an-essay/>

mfence

We've talked about `mfence`.
Used in spinlocks, for instance.

Profiles said: spinlocking didn't take much time.
Empirically: eliminating spinlocks = better than expected!

Exploring the lie

Next step: crate microbenchmarks.

Memory accesses to uncached locations,
or computations,

surrounded by store pairs/mfence/locks.

Use perf to evaluate impact of mfence vs lock.

perf for lock

```
$ perf annotate -s cache_misses
```

```
[...]
0.06 :          4006b0:          and    %rdx,%r10
0.00 :          4006b3:          add     $0x1,%r9
;; random (out of last level cache) read
0.00 :          4006b7:          mov     (%rsi,%r10,8),%rbp
30.37 :         4006bb:          mov     %rcx,%r10
;; foo is cached, to simulate our internal lock
0.12 :          4006be:          mov     %r9,0x200fbb(%rip)
0.00 :          4006c5:          shl     $0x17,%r10
[... Skipping arithmetic with < 1% weight in the profile]
;; locked increment of an in-cache "lock" byte
1.00 :          4006e7:          lock incb 0x200d92(%rip)
21.57 :         4006ee:          add     $0x1,%rax
[...]
;; random out of cache read
0.00 :          400704:          xor     (%rsi,%r10,8),%rbp
21.99 :         400708:          xor     %r9,%r8
[...]
;; locked in-cache decrement
0.00 :          400729:          lock decb 0x200d50(%rip)
18.61 :         400730:          add     $0x1,%rax
[...]
0.92 :          400755:          jne     4006b0 <cache_misses+0x30>
```


lock's effects

Reads take $30 + 22 = 52\%$ of runtime

Locks take $19 + 21 = 40\%$.

perf for mfence

```
$ perf annotate -s cache_misses
[...]  
0.00 :      4006b0:      and    %rdx,%r10  
0.00 :      4006b3:      add     $0x1,%r9  
;; random read  
0.00 :      4006b7:      mov     (%rsi,%r10,8),%rbp  
42.04 :      4006bb:      mov     %rcx,%r10  
;; store to cached memory (lock word)  
0.00 :      4006be:      mov     %r9,0x200fbb(%rip)  
[...]  
0.20 :      4006e7:      mfence  
5.26 :      4006ea:      add     $0x1,%rax  
[...]  
;; random read  
0.19 :      400700:      xor     (%rsi,%r10,8),%rbp  
43.13 :      400704:      xor     %r9,%r8  
[...]  
0.00 :      400725:      mfence  
4.96 :      400728:      add     $0x1,%rax  
0.92 :      40072c:      add     $0x1,%rax  
[...]  
0.36 :      40074d:      jne     4006b0 <cache_misses+0x30>
```

mfence's effects

Looks like the reads take 85% of runtime,
while the mfence takes 15% of runtime.

Bigger picture

Must also look at total # of cycles.

No atomic/fence:	2.81e9 cycles
lock inc/dec:	3.66e9 cycles
mfence:	19.60e9 cycles

That 15% number is a total lie.

Conclusions

- mfence underestimated;
- lock overestimated.

Why?

mfence = pipeline flush,
costs attributed to instructions being flushed.

Summary

Saw a bunch of lies today:

- calling-context lies;
- perf attribution lies.